

Reproducible Research: Peer Assessment 1

Loading and preprocessing the data

We first read the raw data from the .zip file that (we assume) is in our current working directory into the data frame `data`.

```
if (!file.exists("activity.csv")){
  unzip("activity.zip")
}

data <- read.csv("activity.csv")
str(data)
```

```
## 'data.frame': 17568 obs. of 3 variables:
## $ steps : int NA NA NA NA NA NA NA NA NA NA ...
## $ date : Factor w/ 61 levels "2012-10-01","2012-10-02",...: 1 1 1 1 1 1 1 1 1 1 ...
## $ interval: int 0 5 10 15 20 25 30 35 40 45 ...
```

We convert the date column into POSIXct format, which will be useful later when applying the `weekdays()` function. Although irrelevant, we set the time zone to GMT.

```
data$date <- as.POSIXct(data$date, tz="GMT")
str(data)
```

```
## 'data.frame': 17568 obs. of 3 variables:
## $ steps : int NA NA NA NA NA NA NA NA NA NA ...
## $ date : POSIXct, format: "2012-10-01" "2012-10-01" ...
## $ interval: int 0 5 10 15 20 25 30 35 40 45 ...
```

What is mean total number of steps taken per day?

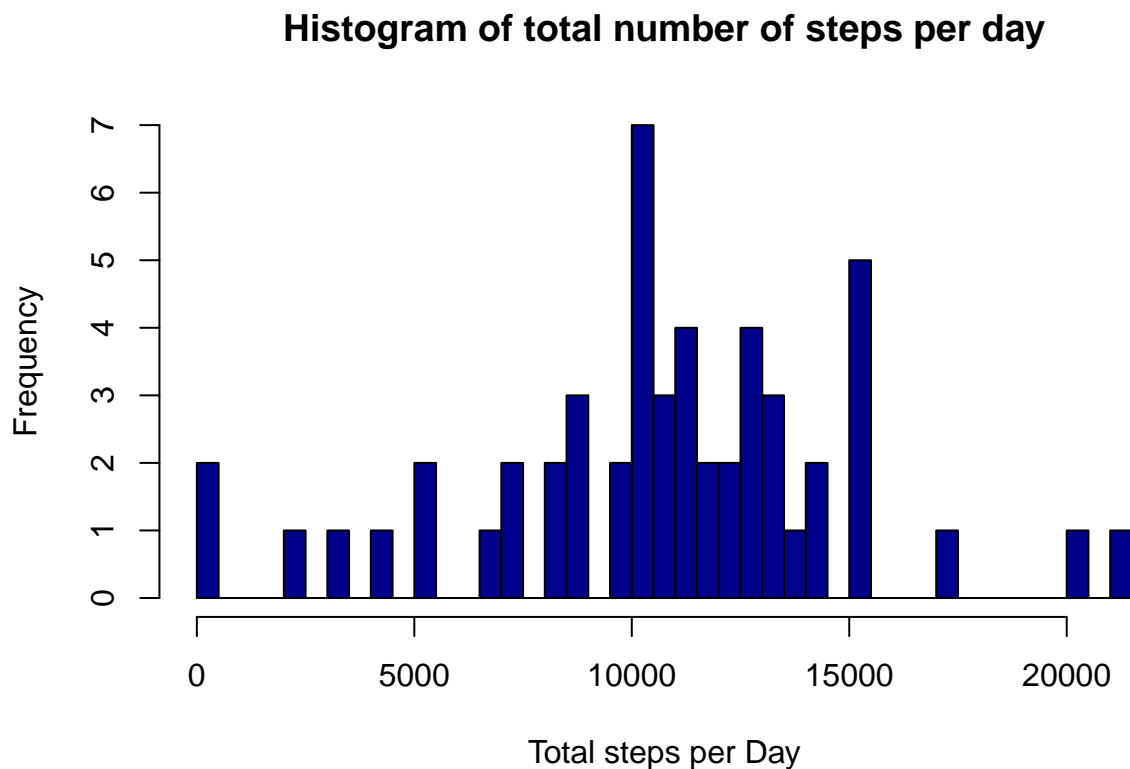
We make a data frame containing the total number of steps for each day. We ignore NA values.

```
totalstepsPerDay <- aggregate(data$steps, by = list(data$date), FUN=sum)
names(totalstepsPerDay) <- c("date", "totalsteps")
head(totalstepsPerDay)
```

```
##      date totalsteps
## 1 2012-10-01         NA
## 2 2012-10-02        126
## 3 2012-10-03       11352
## 4 2012-10-04       12116
## 5 2012-10-05       13294
## 6 2012-10-06       15420
```

We make a histogram of the total number of steps taken each day.

```
hist(totalstepsPerDay$totalsteps, breaks=33, main="Histogram of total number of steps per day", xlab="T
```



Using the total number of steps taken each day, we calculate the mean...

```
mean(totalstepsPerDay$totalsteps, na.rm=TRUE)
```

```
## [1] 10766.19
```

...and the median of these numbers.

```
median(totalstepsPerDay$totalsteps, na.rm=TRUE)
```

```
## [1] 10765
```

What is the average daily activity pattern?

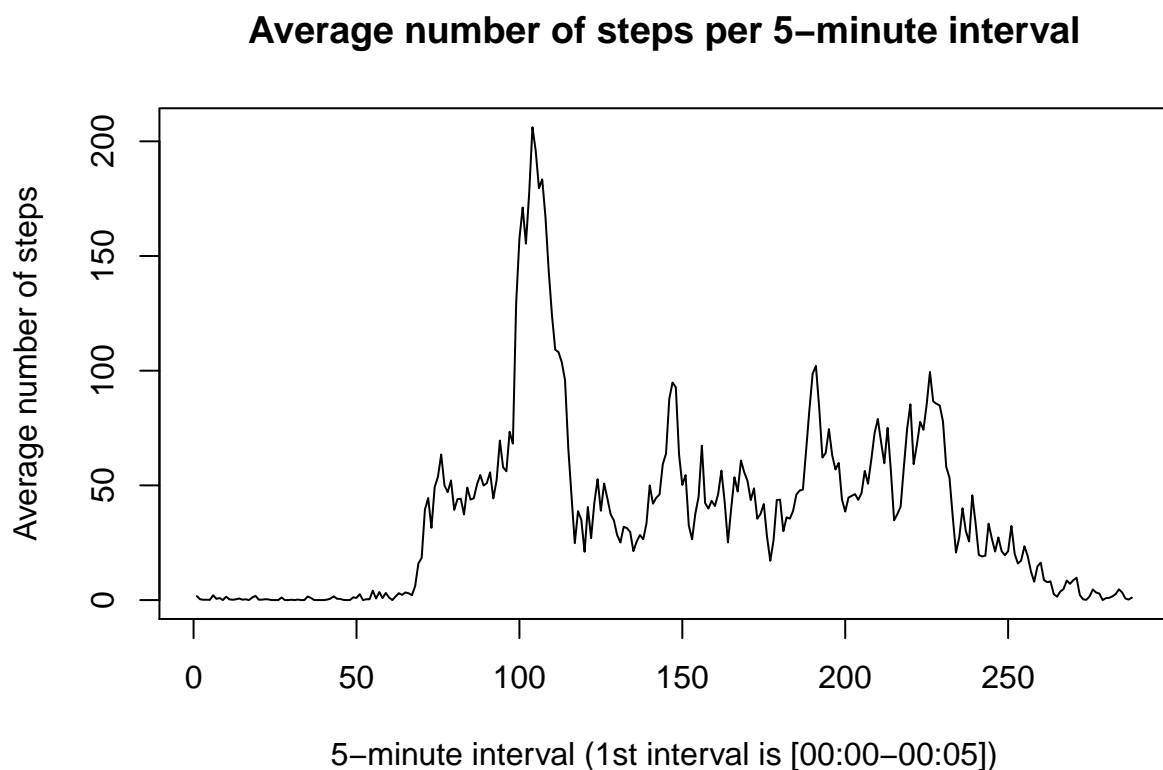
We make a data frame containing the average number of steps taken in every 5-minute interval, averaged across all days and removing NA values.

```
averageStepsPerInterval <- aggregate(data$steps, by = list(data$interval), FUN=mean, na.rm=TRUE)
names(averageStepsPerInterval) <- c("interval", "averagesteps")
head(averageStepsPerInterval)
```

```
##   interval averagesteps
## 1      0      1.7169811
## 2      5      0.3396226
## 3     10      0.1320755
## 4     15      0.1509434
## 5     20      0.0754717
## 6     25      2.0943396
```

We make a time series plot of the 5-minute interval and the average number of steps taken, averaged across all days.

```
plot(seq_along(averageStepsPerInterval$interval), averageStepsPerInterval$averagesteps, type="l", main=
```



We can calculate which interval has the maximum number of steps on average accross all days.

```
which.max(averageStepsPerInterval$averagesteps)
```

```
## [1] 104
```

We can also give the interval in a more proper format.

```
leftIndex<-which.max(averageStepsPerInterval$averagesteps)
rightIndex<-(leftIndex+1) %% dim(averageStepsPerInterval)[1]
maxInterval <- paste("[", floor(averageStepsPerInterval$interval[leftIndex]/100), ":", averageStepsPerInterval$interval[rightIndex]/100, "].")
```

Therefore, the 5-minute interval with the maximum number of steps on average accross all days is [8:35-8:40].

Imputing missing values

We calculate the total number of rows having NA's in the number of steps.

```
sum(is.na(data$steps))
```

```
## [1] 2304
```

We replace each NA value with the average number of steps in the specific 5-minute interval, accross all days, thus creating the data frame `dataNoNA`. Of course, the steps are automatically coerced into numeric (instead of integer).

```
dataNoNA <- data
dataNoNA$steps[is.na(data$steps)] <- rep(averageStepsPerInterval$averagesteps, times = dim(data)[1]/(24*5))
head(dataNoNA)
```

```
##      steps      date interval
## 1 1.7169811 2012-10-01         0
## 2 0.3396226 2012-10-01         5
## 3 0.1320755 2012-10-01        10
## 4 0.1509434 2012-10-01        15
## 5 0.0754717 2012-10-01        20
## 6 2.0943396 2012-10-01        25
```

We perform similar analysis. In particular, we make a data frame containing the total number of steps for each day using `dataNoNA` this time.


```
totalstepsPerDayNoNA <- aggregate(dataNoNA$steps, by = list(dataNoNA$date), FUN=sum)
names(totalstepsPerDayNoNA) <- c("date", "totalsteps")
head(totalstepsPerDayNoNA)
```

```
##      date totalsteps
## 1 2012-10-01  10766.19
## 2 2012-10-02   126.00
## 3 2012-10-03  11352.00
## 4 2012-10-04  12116.00
## 5 2012-10-05  13294.00
## 6 2012-10-06  15420.00
```

We make a histogram of the total number of steps taken each day.

```
hist(totalstepsPerDayNoNA$totalsteps, breaks=33, main="Histogram of total number of steps per day witho
```

Histogram of total number of steps per day without NAs



We can see that this histogram differs from the corresponding histogram for the original dataframe `data` in the y-axis, since there are a lot of NA values taken into account now.

Using the total number of steps taken each day, we calculate the mean...

```
mean(totalstepsPerDayNoNA$totalsteps, na.rm=TRUE)
```

```
## [1] 10766.19
```

...and the median of these numbers.

```
median(totalstepsPerDayNoNA$totalsteps, na.rm=TRUE)
```

```
## [1] 10766.19
```

We can now see that the values for the mean in the processed data frame `dataNoNA` is the same as in the original data frame `data`, and this is because we chose to replace NA values with the mean. On the other hand, the median is different from the case where NA's were not taken into account.

Are there differences in activity patterns between weekdays and weekends?

We create a new factor variable `weekday`, indicating whether a given date is a weekday or weekend day. We use the `dataNoNA` data frame.

```

typeOfDay <-function(date){
  if (weekdays(date) %in% c("Saturday", "Sunday")){
    return("Weekend")
  } else {
    return("Weekday")
  }
}

dataNoNA$weekday <- sapply(dataNoNA$date, typeOfDay)
dataNoNA$weekday <- as.factor(dataNoNA$weekday)
str(dataNoNA)

```

```

## 'data.frame': 17568 obs. of 4 variables:
## $ steps : num 1.717 0.3396 0.1321 0.1509 0.0755 ...
## $ date : POSIXct, format: "2012-10-01" "2012-10-01" ...
## $ interval: int 0 5 10 15 20 25 30 35 40 45 ...
## $ weekday : Factor w/ 2 levels "Weekday","Weekend": 1 1 1 1 1 1 1 1 1 1 ...

```

We now find the mean number of steps accross all weekdays and accross all weekends.

```

averageStepsPerIntervalAndWeekdayNoNA <- aggregate(data$steps, by = list(dataNoNA$interval, dataNoNA$weekday),
names(averageStepsPerIntervalAndWeekdayNoNA) <- c("interval", "weekday", "averagesteps")
head(averageStepsPerIntervalAndWeekdayNoNA)

```

```

## interval weekday averagesteps
## 1 0 Weekday 2.333333
## 2 5 Weekday 0.4615385
## 3 10 Weekday 0.1794872
## 4 15 Weekday 0.2051282
## 5 20 Weekday 0.1025641
## 6 25 Weekday 1.5128205

```

We make a panel plot containing a time series plot of the 5-minute interval (x-axis) and the average number of steps taken, averaged across all weekday days or weekend days (y-axis).

```

library(lattice)
xyplot(averageStepsPerIntervalAndWeekdayNoNA$averagesteps ~ averageStepsPerIntervalAndWeekdayNoNA$interval)

```

Average number of steps per 5-minute interval

